

Provider Aliases — User Guide

claude-providers turns every LLM API key you keep in your keys file into its own Claude Code alias, pointed at that provider's strongest model. It reuses the multi-account toolkit's machinery (shared plugins/history, automatic `.bashrc/.zshrc` wiring) so a provider alias behaves like a normal `claudeN` account, just talking to a different model backend.

It is **fully dynamic**: the provider list, base URLs, transports, and model IDs are all derived at run time from your keys file + the [models.dev](#) catalog. Nothing is hardcoded. Re-run it any time to pick up new keys or refresh existing providers.

1. Prerequisites

Need	Why	Install
The toolkit installed	provides <code>lib.sh</code> , the alias file, <code>sync-state</code>	<code>bash scripts/install.sh</code>
A keys file	one <code>export NAME_API_KEY=...</code> per provider	<code>~/api_keys.sh</code> (default) or <code>pass --keys-file</code>
<code>jq</code> , <code>python3</code> , <code>curl</code>	catalog parse + resolve + fetch	system package manager
<code>claude-code-router</code> (<i>only for routed providers</i>)	translate Anthropic↔OpenAI/Gemini	<code>npm install -g @musistudio/claude-code-router</code>
<code>submodules/LLMsVerifier</code> built (<i>optional</i>)	full “does this model work?” verification	see §7

Your keys file is the **only** place secrets live. `claude-providers` reads the variable *names* from it (without executing it) to discover providers, and the launch wrapper reads the *value* at the moment you start a session. No key is ever written into the repo, the alias file, or the per-provider env files.

2. Quick start

```
claude-providers sync          # discover + create/refresh all
                               provider aliases
source ~/.local/share/claude-multi-account/aliases.sh  # or open
                               a new shell
claude-providers list          # see what you got
deepseek                      # launch Claude Code on DeepSeek's
                               strongest model
```

3. Commands

Command	Does
claude-providers sync <i>(default)</i>	Discover every LLM key, resolve it via models.dev, verify it with live sentinel + tool-calling probes (see §7), and create/refresh one alias per provider. Idempotent — safe to run repeatedly.
claude-providers list	Table of verified provider aliases only (safe to launch): alias name, provider, transport, strong + fast model.
claude-providers list-all	Every installed provider alias, any status.
claude-providers list-faulty	Only aliases with an issue (failed / unverified / pending) — the filtered-out set.
claude-providers show <id>	Print the resolved env file for one provider.
claude-providers remove <id>	Remove the alias + env file; back up and remove the config dir.
claude-providers add --from-key VAR --id PROVIDER	Register a key→provider mapping (when a key's name doesn't match models.dev), then sync.

Useful flags

- `--keys-file PATH` — use a different keys file (default `~/api_keys.sh`).
- `--no-verify` — skip verification; create aliases regardless.
- `--offline` — don't fetch `models.dev`; use the local cache (errors if none).
- `--dry-run` — print what would change; write nothing.

4. How a provider alias is named and what it runs

- **Alias name** defaults to the `models.dev` provider id (`deepseek`, `groq`, `mistral`, ...). Want something shorter like `dseek`? Set it in `scripts/providers/overrides.json` (see §6).
- **Strong model** → Claude Code's main model (`ANTHROPIC_MODEL`). Chosen as the provider's most capable: reasoning-capable first, then newest, then largest context.
- **Fast model** → Claude Code's background model (`ANTHROPIC_SMALL_FAST_MODEL`), chosen as the cheapest tool-call-capable model.
- **Transport:**
 - `native` — provider speaks the Anthropic API directly (`models.dev` npm is `@ai-sdk/anthropic`). The alias runs `claude` with `ANTHROPIC_BASE_URL` / `ANTHROPIC_AUTH_TOKEN` / `ANTHROPIC_MODEL` set.
 - `router` — provider is OpenAI-compatible / Gemini. The alias launches through `claude-code-router` (`ccr` code), which translates the protocol.

5. Config dirs, plugins, and shared state

Each provider gets `~/.claude-prov-<id>`, which symlinks the same shared items as your accounts (`projects`, `todos`, `plugins`, `settings.json`, `CLAUDE.md`, ...). So **every plugin you have in `claude1..N` is available in every provider session**, and history/projects are unified. `sync` also force-enables the always-on set in shared settings: **superpowers, systematic-debugging, frontend-design, code-review**.

Provider dirs are deliberately **excluded from account auto-detection**, so they never get merged into your real accounts' auth/identity and never interfere with `claude-unify` or `claude-add-account`. Your existing `claudeN` aliases keep working exactly as before, and you can still create new accounts the same way.

Cross-alias session visibility (v1.5.0+)

Sessions created under **any** alias — `claude1`, `claude2`, `deepseek`, `opencode`, `xiaomi`, or any other — are visible from **every** other alias via `/resume`. This works automatically:

1. **Before launch:** `claude-sync-state pull` merges every account's and provider's `.claude.json` into the launching dir. This includes `lastSessionId` (what `/resume` uses), `allowedTools`, MCP config, and all other project settings.
2. **After exit:** `claude-sync-state push` merges the post-session state back out, so the next alias to launch picks up the new session.
3. **Session files** are in `~/.claude-shared/sessions/`, which every alias symlinks to — so the session data itself was always shared; the new work ensures `.claude.json` project metadata is also merged.

Example workflow:

```
# Start a project as deepseek
cd /path/to/my-project
deepseek          # bare launch: auto-creates the project's
                  session
# ... work for a while, then exit

# Continue the same project as opencode
opencode          # bare launch: AUTO-RESUMES the same project
                  session – no /resume needed
```

A **bare** alias launch (no arguments) now auto-resumes — or first-time creates — the one long-lived session keyed to the project root, so switching aliases continues the same conversation automatically (see §12 of `../Claude_Multi_Account_Fine_Tuning.md` and `SESSION_COLOR.md`). You only need `/resume` to switch to a **different** session than the project's own.

Performance: adds ~1-2 seconds per launch (jq merge across all dirs). Same overhead that `claudeN` aliases already have.

Note: provider dirs are still excluded from `cma_detect_accounts` (used by `claude-unify` and `claude-add-account`). Only `claude-sync-state` includes them, so the existing account management is unaffected.

Multi-alias system (v1.6.0+)

By default, sync creates **one alias per provider** using the 2 strongest models. With `--multi`, it verifies **ALL models** for each provider and creates **multiple aliases** — provider, provider2, provider3, etc. — each with a different pair of strong + fast models.

Standard: 1 alias per provider (2 models)

```
claude-providers sync
```

Multi: verify all models, create multiple aliases

```
claude-providers sync --multi
```

Options

```
claude-providers sync --multi --max-aliases 10 --min-score 20
```

How verification works (v1.14.0+):

Each model is tested via live HTTP probes against the provider's chat endpoint: a sentinel probe (sends "Reply with exactly: VERIFY_OK" and **requires** the response to contain VERIFY_OK) and a tool-calling probe (requires the model to actually emit a tool call). Models are scored on 7 dimensions (0-100):

Dimension	Weight	What it checks
Existence + valid response	25pts	Model exists, returns real content
Tool calling	20pts	Supports function/tool_calls
Reasoning	15pts	Has chain-of-thought / reasoning_content
Context window	15pts	≥8K tokens (log scale)
Streaming	10pts	SSE streaming support
Latency	10pts	Response under 2s (full) or 5s (half)
Free tier	5pts	\$0 input cost

Tool calling is a hard gate, not just score. A model that answers chat but never calls tools is marked **not verified**, no matter how high its score — Claude Code is entirely tool-driven, so a tool-less model always breaks at runtime.

Anti-bluff detection prevents false positives: - HTTP 200 with error body (JSON error wrapped in success) - Empty response content - **200 responses that lack the VERIFY_OK sentinel** (proxy fallback, silent model swap, canned text) - Boilerplate "I'm unable to" responses - Models that claim capability but don't deliver

Alias pairing: - Models sorted by score, paired 2 per alias (strong + fast) - Odd count: last model used for both positions in last alias - Single model: used for both positions - Default max: 5 aliases per provider (configurable)

Verification cache: results cached for 24h to avoid re-testing on every sync. Cache stored in `~/.local/share/claude-multi-account/providers/verification_cache.json`. The cache carries a schema version; results written by older verification logic (e.g. before tool calling was required) are ignored automatically.

6. Overrides — `scripts/providers/overrides.json`

Per-provider pins. Empty by default. Any field you set wins over the `models.dev`-derived value:

```
{
  "deepseek": {
    "alias": "dseek",
    "transport": "native",
    "base_url": "https://api.deepseek.com/anthropic",
    "strong_model": "deepseek-reasoner",
    "fast_model": "deepseek-chat"
  }
}
```

This is how you (for example) promote DeepSeek to its native Anthropic endpoint (`/anthropic`) instead of routing it, or give it the short alias `dseek` — all without touching any code.

`scripts/providers/key-aliases.json`

Maps a key variable name to a `models.dev` provider id, for the cases where the name differs (e.g. `CODESTRAL_API_KEY` → `mistral`, `ZAI_API_KEY` → `zai`). Use `claude-providers add --from-key VAR --id PROVIDER` to append entries.

7. Verification — what sync actually checks (v1.14.0+)

By default sync verifies every provider with **two live probes** against its chat endpoint, using the exact model the alias will run:

1. **Sentinel probe** — asks the model to reply with exactly `VERIFY_OK` and requires that token in the response. A 200 without it (or with an error object smuggled into the body) is a

bluff: the endpoint answered *something*, not the requested model.

2. **Tool-calling probe** — requires the model to emit a real tool call (`tool_calls` / `tool_use`). Claude Code is tool-driven; a model that only chats is useless in practice.

Verdicts:

- **verified** — both probes passed. The alias is created and launchable.
- **failed** — a definitive problem: auth/billing rejection (401/402/403), model missing (404), missing sentinel, error-in-200, or no tool call. The alias is **not activated** (and the launch gate refuses to run it).
- **unverified** — the probes were inconclusive (rate-limit 429, 5xx, timeout, no network, no key). The alias is still created but the **launch gate** refuses to start it until a later sync reaches a verdict (override with `--force` — not recommended).

This replaced the old best-effort check (a bare `GET /models`), which proved only that the key was accepted — not that the model can actually run Claude Code. Both Anthropic-native (`/v1/messages`) and OpenAI-compatible (`/chat/completions`) endpoint shapes are probed in their native format.

For an additional authoritative layer (“can this model genuinely see and describe my code?”) build the `LLMsVerifier` submodule:

```
git submodule update --init --depth 1 submodules/LLMsVerifier
cd submodules/LLMsVerifier && go build -o bin/model-
verification ./llm-verifier/cmd/model-verification
```

Once `submodules/LLMsVerifier/bin/model-verification` exists, `sync` prefers it as the existence check. `sync` also runs the two-round semantic code-visibility check (sentinel + independent judge) when it is available; billing/auth rejections there demote the alias, while transient judge/infra errors are an honest `SKIP` that never demotes. (Requires the Go toolchain; the build reads keys from your keys file.)

8. Known limitation — default session color

The original goal was for each provider session to default its `/color` to **purple**. Investigation of the installed Claude Code (2.1.195) found that `/color` is **session-scoped and TUI-only — it is never persisted to disk**, and there is no settings key, hook, or environment variable to set a default color. So this cannot be automated with the current Claude Code.

Workaround: type `/color purple` at the start of a provider session. `purple` is a valid color. If a future Claude Code adds a persistable color setting, `sync` will seed it automatically.

9. Troubleshooting

Symptom	Cause / fix
provider 'X' needs <code>claude-code-router</code>	Routed provider, <code>ccr</code> not installed → <code>npm install -g @musistudio/claude-code-router</code> .
<code>\$NAME_API_KEY</code> is empty	The key isn't set in your keys file.
offline and no <code>models.dev</code> cache	Run <code>claude-providers sync</code> once online to populate the cache.
A key didn't become an alias	It's unmapped (no <code>models.dev</code> match) or classified <code>vcs/infra</code> . Map it with <code>add</code> , or check <code>scripts/providers/evidence/mapping-report.md</code> .
Alias missing after <code>sync</code>	<code>source ~/.local/share/claude-multi-account/aliases.sh</code> or open a new shell.

10. Remote distributed deployment & heavy testing

For heavy testing that depends on **real production services**, the toolkit can boot the full LLMsVerifier System on a remote host (registered in `config/containers/nezha.env`) via the `containers` submodule + `podman`, and run real per-provider verification through it.

- **Deployment guide + boot procedure:** `config/containers/llmsverifier/README.md` (services, ports, the exact `podman-compose` boot steps, and every root-caused deployment fix).
- **What runs:** the `llm-verifier` API + an observability tier (`prometheus` + `grafana` with an auto-provisioned datasource & dashboard + `node-exporter`).
- **Heavy testing:** the bundled model-verification tool runs a real “Do you see my code?” check against each provider’s live API:

```
ssh <user>@<host> 'podman run --rm --env-file ~/helix-system/
    llmsverifier/.env \
    llm-verifier-mv:nezha --provider deepseek --model deepseek-
    chat --verbose'
```


- **Evidence:** every verification + deployment claim is captured under docs/qa/20260616-infra/ (proofs, sweeps, security, observability).

11. Command quick-reference

```
# Provider aliases
claude-providers sync           # discover keys -> create/
    refresh aliases
claude-providers sync --no-verify # skip verification (faster)
claude-providers list           # alias | provider |
    transport | models
claude-providers show <id>      # one provider's resolved
    env
claude-providers remove <id>    # remove alias + env, back
    up config dir
claude-providers add --from-key VAR --id PROVIDER # register
    mapping + sync

# Launch a provider session (after: source the alias file or open
    a new shell)
deepseek                        # native provider -> claude
<router-provider>              # routed provider -> claude
    via ccr
deepseek -p "your prompt"      # non-interactive print mode

# Accounts (unchanged, still works)
claude-add-account --alias claudeN # add a Claude account
claude-list-accounts              # status of all accounts
claude-unify                      # re-merge shared state

# Docs
claude-export-docs                # regenerate HTML/PDF/DOCX
```

12. Individual provider notes

z.ai Coding Plan (zai-coding-plan)

The [z.ai](#) Coding Max-Yearly Plan provides access to Zhipu AI's flagship GLM models through a dedicated coding-optimized API endpoint. This is separate from the base z.ai plan — the Coding Plan endpoint uses `api.z.ai/api/coding/paas/v4` and includes **Free access** to glm-5.2 (the flagship with 1M context and reasoning) and glm-4.7 (the fast model with 204k context, reasoning, and tool_call support), among others.

How the alias works

The key variable `ZAI_API_KEY` in your keys file is mapped to the `zai-coding-plan` provider ID via `scripts/providers/key-aliases.json`. An override in `scripts/providers/overrides.json` pins the strong model to `glm-5.2` and the fast model to `glm-4.7`, which are the optimal choices for coding workloads on this plan.

Transport is **router** — the alias launches through `claude-code-router` (`ccr` code), which translates the Anthropic protocol to the OpenAI-compatible `z.ai` API.

Models available on the Coding Plan

All models are free on the Coding Max-Yearly Plan:

Model	Context	Reasoning	Tool Call	Notes
glm-5.2	1M tokens	Yes	Yes	Flagship — newest, most capable
glm-5.1	—	Yes	Yes	Intermediate
glm-5-turbo	—	—	—	Turbo variant of glm-5
glm-5	—	—	—	Base generation 5
glm-4.7	204k tokens	Yes	Yes	Fast model — reasoning + tool_call
glm-4.6	—	—	—	Mid-range
glm-4.5-air	—	—	—	Lightweight
glm-4.5	—	—	—	Base model

Note on glm-5.2 rate limits: The Coding Max-Yearly Plan enforces a Fair Usage Policy on the flagship model: rapid-fire requests (e.g., repeated programmatic calls at short intervals) may trigger a temporary rate limit (error code 1313). Normal interactive use via `zai-coding-plan` at human pace is unaffected.

Setup

Everything is already configured on this host, but for a fresh install:

1. Ensure `ZAI_API_KEY` is exported in your keys file (`~/api_keys.sh`).

2. Run `claude-providers sync` to discover the key and create the alias.
3. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
4. Run `zai-coding-plan` to start a Claude Code session on glm-5.2.

No additional manual steps are needed — the key alias, model override, and endpoint are all defined in the provider config files.

Usage example

```
zai-coding-plan                                # launch Claude Code on glm-5.2
zai-coding-plan -p "explain this function"      # non-interactive
print mode
```

Xiaomi MiMo (xiaomi)

Xiaomi MiMo is Xiaomi's LLM platform. Unlike most OpenAI-compatible providers in this toolkit, MiMo exposes a genuine **Anthropic-native endpoint** (<https://api.xiaomimimo.com/anthropic>, POST `/anthropic/v1/messages`) that accepts `Authorization: Bearer <key>` and returns native Anthropic-format responses. Because of this, the `xiaomi` alias uses **native transport** — Claude Code talks to MiMo directly with no `claude-code-router` (`ccr`) dependency, the same way the `deepseek` alias works.

How the alias works

The key variable `XIAOMI_MIMO_API_KEY` in your keys file does **not** match the `models.dev` `xiaomi` provider's documented env name (`XIAOMI_API_KEY`), so `scripts/providers/key-aliases.json` maps `XIAOMI_MIMO_API_KEY` → `xiaomi`. An override in `scripts/providers/overrides.json` pins:

- **transport** `native`
- **base_url** `https://api.xiaomimimo.com/anthropic` (the Anthropic-native endpoint, overriding the catalog's OpenAI-compatible `/v1` URL)
- **strong_model** `mimo-v2.5-pro`
- **fast_model** `mimo-v2-flash`

The pinning is deliberate: `models.dev` lists a `mimo-v2.5-pro-ultraspeed` id that the **live API does not serve**, so the override guarantees only live-served model ids are used.

Models

MiMo's text-generation models (all support tool/function calling and reasoning; verified live 2026-06-19):

Model	Context	Reasoning	Tool Call	Notes
mimo-v2.5-pro	1M tokens	Yes	Yes	Flagship — strongest (alias default)
mimo-v2.5	1M tokens	Yes	Yes	Omni / multimodal (image, audio, video)
mimo-v2-pro	—	Yes	Yes	Legacy v2 Pro
mimo-v2-omni	256k tokens	Yes	Yes	Omni / multimodal (v2)
mimo-v2-flash	256k tokens	Yes	Yes	Fast model — cheapest tier (alias fast)

Non-text models exist (mimo-v2.5-asr, mimo-v2.5-tts*, mimo-v2-tts) but are speech-recognition / text-to-speech and cannot serve chat or code, so they are intentionally not wired as aliases.

Verified live

- GET /v1/models → HTTP 200, 10 models served (the 5 text models above + 5 ASR/TTS).
- POST /anthropic/v1/messages with Authorization: Bearer → HTTP 200, native Anthropic response (type:"message", content:[{type:"text"},{type:"thinking"}]) for both mimo-v2.5-pro and mimo-v2-flash.
- POST /v1/chat/completions with tools[] → finish_reason:"tool_calls", valid tool-call array, reasoning_content present (tool calling works).
- Streaming (stream:true) → SSE chat.completion.chunk deltas.
- Rate limits: 100 RPM / 10M TPM per account for text models (per-account, not per-key); expect 429 under load, retry with backoff.

Setup

Everything is already configured, but for a fresh install:

1. Ensure XIAOMI_MIMO_API_KEY is exported in your keys file (~/.api_keys.sh).
2. Run `claude-providers sync` to discover the key and create the alias.

3. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
4. Run `xiaomi` to start a Claude Code session on `mimo-v2.5-pro`.

Usage example

```
xiaomi                                # launch Claude Code on mimo-  
                                     v2.5-pro  
xiaomi -p "explain this function"    # non-interactive print mode
```

OpenCode Zen (opencode)

OpenCode Zen is OpenCode's curated AI gateway — a tested and verified list of models from multiple providers, accessed through a single API key. Zen includes **21 free models** (all \$0 cost, all support tool calling and reasoning), plus 49 paid models from OpenAI, Anthropic, Google, DeepSeek, and others.

The free models are available for a limited time while OpenCode collects feedback. One of them — **Big Pickle** — is a stealth model (its true identity is not disclosed; observed as `deepseek-v4-flash` behind the scenes).

How the alias works

The key variable `ZEN_API_KEY` (or `ApiKey_Opencode_Zen`) in your keys file is mapped to the `opencode` provider ID via `scripts/providers/key-aliases.json`. An override in `scripts/providers/overrides.json` pins:

- **strong_model** `big-pickle` — free stealth model, 200K context, reasoning + tool_call
- **fast_model** `deepseek-v4-flash-free` — free, 200K context, reasoning + tool_call

Transport is **router** — the alias launches through `claude-code-router` (ccr code), which translates the Anthropic protocol to the OpenAI-compatible Zen API (<https://opencode.ai/zen/v1/chat/completions>).

No transport or `base_url` override is needed — the `models.dev` catalog already has the correct values (`@ai-sdk/openai-compatible` → `router`, <https://opencode.ai/zen/v1>).

Free models available on Zen

All free models support tool calling and reasoning (verified live 2026-06-20):

Model	Context	Notes
big-pickle	200K	Stealth model — alias default (strong)
deepseek-v4-flash-free	200K	Alias fast model
mimo-v2.5-free	200K	Xiaomi MiMo
mimo-v2-pro-free	1M	Xiaomi MiMo Pro
mimo-v2-flash-free	262K	Xiaomi MiMo Flash
mimo-v2-omni-free	262K	Xiaomi MiMo Omni
nemotron-3-ultra-free	1M	NVIDIA Nemotron
nemotron-3-super-free	204K	NVIDIA Nemotron Super
north-mini-code-free	256K	North Mini Code
glm-4.7-free	204K	Zhipu GLM
glm-5-free	204K	Zhipu GLM 5
grok-code	256K	xAI Grok Code
kimi-k2.5-free	262K	Moonshot Kimi
minimax-m2.1-free	204K	MiniMax
minimax-m2.5-free	204K	MiniMax
minimax-m3-free	200K	MiniMax
qwen3.6-plus-free	262K	Alibaba Qwen
ring-2.6-1t-free	262K	Ring
hy3-preview-free	256K	HY3 Preview
ling-2.6-flash-free	262K	Ling (no reasoning)
trinity-large-preview-free	131K	Trinity (no reasoning)

Paid models start at \$0.05/1M tokens (GPT-5 Nano) and go up to \$30/1M (GPT-5.5 Pro). Claude models are available from \$1/1M (Haiku 4.5) to \$15/1M (Opus 4.1).

Verified live

- GET /v1/models → HTTP 200 (model list returned).
- POST /v1/chat/completions with big-pickle → HTTP 200, correct text response, cost=\$0, reasoning_content present. Stealth alias observed as deepseek-v4-flash.
- POST /v1/chat/completions with deepseek-v4-flash-free → HTTP 200, correct text, cost=\$0.
- Additional free models (mimo-v2.5-free, nemotron-3-ultra-free, north-mini-code-free) → all HTTP 200, all cost=\$0.
- Streaming (stream:true) → SSE chat.completion.chunk deltas.

Setup

Everything is already configured, but for a fresh install:

1. Ensure ZEN_API_KEY (or ApiKey_Opencode_Zen) is exported in your keys file (~/.api_keys.sh).
2. Run `claude-providers sync` to discover the key and create the alias.
3. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
4. Run `opencode` to start a Claude Code session on big-pickle (via ccr).

Usage example

```
opencode                                # launch Claude Code on big-
    pickle (via ccr)
opencode -p "explain this function" # non-interactive print mode
```

Notes

- The alias uses **router transport** (ccr) because Zen's free models use OpenAI-compatible format, not Anthropic native format.
- Big Pickle is a stealth model — the actual model served may vary. This is by design per OpenCode's documentation.
- To use paid Claude models on Zen instead of free models, override `strong_model` and `fast_model` in `scripts/providers/overrides.json` (e.g. `claude-haiku-4.5` and `claude-sonnet-4-6`), and set `transport` to `native` with `base_url` to `https://opencode.ai/zen/v1/messages`.

Chutes (chutes)

Chutes is a pay-per-use AI inference platform with **TEE (Trusted Execution Environment)** models — all models run in secure enclaves for data privacy. The API is OpenAI-compatible at <https://llm.chutes.ai/v1>.

How the alias works

The key variable `CHUTES_API_KEY` in your keys file maps to the chutes provider ID. An override in `scripts/providers/overrides.json` pins:

- **strong_model** `zai-org/GLM-5.2-TEE` — newest GLM model, 202K context
- **fast_model** `Qwen/Qwen3.6-27B-TEE` — fast Qwen model, 262K context

Transport is **router** — the alias launches through claude-code-router (ccr code), which translates the Anthropic protocol to the OpenAI-compatible Chutes API.

Models available on Chutes

All models are TEE-enabled (Trusted Execution Environment):

Model	Context	Notes
zai-org/ GLM-5.2-TEE	202K	Newest GLM — alias default (strong)
zai-org/GLM-5.1-TEE	202K	GLM 5.1
zai-org/GLM-5-TEE	202K	GLM 5
Qwen/ Qwen3.5-397B-A17B-TEE	262K	Largest Qwen model
Qwen/ Qwen3.6-27B-TEE	262K	Alias fast model
Qwen/ Qwen3-32B-TEE	40K	Qwen 32B
Qwen/ Qwen3-235B-A22B-Thinking-2507-TEE	262K	Thinking/reasoning model
deepseek-ai/ DeepSeek-V3.2-TEE	131K	DeepSeek V3.2
moonshotai/Kimi-K2.5-TEE	262K	Kimi K2.5
moonshotai/Kimi-K2.6-TEE	262K	Kimi K2.6
MiniMaxAI/ MiniMax-M2.5-TEE	196K	MiniMax M2.5
google/ gemma-4-31B-turbo-TEE	131K	Google Gemma 4
unsloth/Mistral-Nemo-Instruct-2407-TEE	131K	Mistral Nemo

Important: Chutes is **pay-per-use** — the account must be funded to use models. Add balance at <https://chutes.ai>. The API key and configuration are correct, but requests will return “Quota exceeded” if the account balance is \$0.

Setup

1. Ensure CHUTES_API_KEY is exported in your keys file (~/.api_keys.sh).
2. Fund your Chutes account at <https://chutes.ai>.
3. Run `claude-providers sync` to discover the key and create the alias.
4. `source ~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
5. Run `chutes` to start a Claude Code session on `zai-org/GLM-5.2-TEE`.

Usage example

```
chutes                                # launch Claude Code on
                                     GLM-5.2-TEE (via ccr)
chutes -p "explain this function"    # non-interactive print mode
```

Verified

- API endpoint <https://llm.chutes.ai/v1/chat/completions> responds correctly
- Authentication with Authorization: Bearer `cpk_...` works
- All 13 TEE models are accessible (require funded account)
- OpenAI-compatible format confirmed

Poe (poe)

Poe is a universal AI platform with **382 models** from all major providers — chat, code, image generation, video generation, TTS, and STT. The API is OpenAI-compatible at <https://api.poe.com/v1>.

How the alias works

The key variable `POE_API_KEY` (or `ApiKey_Poe`) in your keys file maps to the poe provider ID. Transport is **router** — the alias launches through `claude-code-router` (`ccr code`).

Aliases

Alias	Strong Model	Fast Model	Focus
poe	claude-sonnet-4.6	gpt-5.4-mini	Primary — best Claude

Alias	Strong Model	Fast Model	Focus
poe2	gpt-5.5	deepseek-v4-pro-e	GPT-focused
poe3	grok-4	gemini-3.1-pro	Alternative providers

Model categories (382 total)

Category	Count	Examples
Chat/Reasoning	130	claude-opus-4.8, gpt-5.5, grok-4, deepseek-v4-pro, gemini-3.1-pro, qwen3.7-max
Code	16	claude-code, gpt-5.3-codex, qwen3-coder-next, kimi-k2.7-code, seed-2.0-code
Image Generation	40	flux-2-pro, imagen-4, stable-diffusion3.5, qwen-image-2, grok-imagine-image
Video Generation	17	sora-2-pro, veo-3.1, kling-3.0, runway-gen-4.5, pixverse-v5.6
TTS/Voice	12	elevenlabs-v3, gemini-3.1-flash-tts, cartesia-ink-whisper, orpheus-tts
STT/Speech	1	whisper-v3-large-t
Other	166	perplexity-search, exa-research, hailuo, minimax-m3, glm-5.2, and more

Supported capabilities

- **Chat completions** — OpenAI-compatible /v1/chat/completions
- **Tool calling** — verified on claude-sonnet-4.6, gpt-5.4-mini, deepseek-v4-pro-e, grok-4
- **Streaming** — standard SSE streaming supported
- **Reasoning** — reasoning_content field in responses
- **Image input** — vision models support image URLs
- **Video generation** — via separate video endpoint
- **TTS/STT** — text-to-speech and speech-to-text models

Setup

1. Ensure POE_API_KEY is exported in your keys file (~/.api_keys.sh).
2. Run `claude-providers sync` to discover the key and create the alias.
3. source `~/.local/share/claude-multi-account/aliases.sh` (or open a new shell).
4. Run `poe` to start a Claude Code session on `claude-sonnet-4.6` (via `ccr`).

Usage example

```
poe                                # launch Claude Code on
                                claude-sonnet-4.6 (via ccr)
poe2                              # launch on gpt-5.5 (via ccr)
poe3                              # launch on grok-4 (via ccr)
poe -p "explain this function"    # non-interactive print mode
```

Verified

- API endpoint `https://api.poe.com/v1/chat/completions` responds correctly
- Authentication with Authorization: Bearer `sk-poe-...` works
- Tool calling verified on multiple models
- 382 models accessible across all categories
- OpenAI-compatible format confirmed